

Language Modeling


Natural Language Processing

Emory University

Jinho D. Choi



Probability



Sunny

5 days



Cloudy

3 days




Snowy

2 days

Probability of tomorrow being **cloudy**?

$$P(\text{cloudy}) = \frac{C(\text{cloudy})}{C(\text{sunny}) + C(\text{cloudy}) + C(\text{snowy})} = \frac{3}{10}$$


Conditional Probability



Probability of tomorrow being **cloudy** if today is **snowy**?

$$P(\text{cloudy}|\text{snowy}) = \frac{C(\text{snowy}, \text{cloudy})}{C(\text{snowy})} = \frac{1}{2}$$


Conditional Probability



Probability of tomorrow being **cloudy** if
today and yesterday are **snowy**?

$$P(\text{cloudy}|\text{snowy}, \text{snowy}) = \frac{C(\text{snowy}, \text{snowy}, \text{cloudy})}{C(\text{snowy}, \text{snowy})} = \frac{1}{1} = 1$$


Joint Probability



Probability of 2 consecutive days being **cloudy, sunny**?

$$P(\text{cloudy}, \text{sunny}) = P(\text{cloudy}) \cdot P(\text{sunny}|\text{cloudy})$$

Joint Probability



Probability of next 3 days being **snowy, cloudy, sunny**?

$$P(\text{snowy}, \text{cloudy}, \text{sunny}) = P(\text{snowy}) \cdot P(\text{cloudy} | \text{snowy}) \cdot P(\text{sunny} | \text{snowy}, \text{cloudy})$$

N-gram Models

1-gram (Unigram)

$$P(w_i) = \frac{C(w_i)}{\sum_{\forall k} C(w_k)} = \frac{C(w_i)}{\boxed{N}}$$

of tokens

token
vs
type?

2-gram (Bigram)

$$P(w_{i+1}|w_i) = \frac{C(w_i, w_{i+1})}{\sum_{\forall k} C(w_i, w_k)} = \frac{C(w_i, w_{i+1})}{C(w_i)}$$

N-gram Models

- Unigram model
 - Given any word w , it shows how **likely** w appears in context.
 - This is known as the **likelihood** (probability) of w , written as $P(w)$.
 - How likely does the word “**Emory**” appear in context?

Emory University was found as **Emory** College by John **Emory** .
Emory University is 16th among the colleges and universities in US .

$$P(\textit{Emory}) = \frac{4}{23} \approx 0.1739$$

- Does this mean “**Emory**” appears 17.39% time in any context?
- How can we measure more accurate **likelihoods**?



N-gram Models

- Bigram model
 - Given any words w_i and w_j in sequence, it shows the **likelihood** of w_j following w_i in context.
 - This can be represented as the conditional probability of $P(w_j|w_i)$.
 - What is the most likely word following “Emory”?

Emory University was found as Emory College by John Emory.
Emory University is the 20th among the national universities in US.

$$\arg \max_k P(w_k | Emory)$$

$P(University Emory)$	$=$	$\frac{2}{4}$	$= 0.5$
$P(College Emory)$	$=$	$\frac{1}{4}$	$= 0.25$
$P(. Emory)$	$=$	$\frac{1}{4}$	$= 0.25$



Maximum Likelihood

$$x_1^n = x_1, \dots, x_n$$

Chain rule

$$P(x_1^n) = P(x_1) \cdot P(x_2|x_1) \cdot P(x_3|x_1^2) \cdots P(x_n|x_1^{n-1})$$

Any practical issue?

(x_1, \dots, x_k) can be very sparse.

Markov assumption

$$P(x_k|x_1^{k-1}) \approx P(x_k|x_{k-1})$$

$$P(x_1^n) \approx P(x_1) \cdot P(x_2|x_1) \cdot P(x_3|x_2) \cdots P(x_n|x_{n-1})$$



Maximum Likelihood

- Maximum likelihood
 - Given any word sequence w_i, \dots, w_n , how **likely** this sequence appears in context.
 - This can be represented as the joint probability of $P(w_j, \dots, w_n)$.
 - How likely does the sequence “you know” appears in context?

you know , I know you know that you do .

$$P(\text{you}, \text{know}) = \frac{2}{11} \text{ not } 10?$$

Chain rule

$$P(\text{you}) \cdot P(\text{know}|\text{you}) = \frac{3}{11} \cdot \frac{2}{3} = \frac{2}{11}$$



Maximum Likelihood

$$P(w_1, w_2, \dots, w_n) = P(w_1^n)$$

$$P(w_1) \cdot P(w_2|w_1) \cdot P(w_3|w_1^2) \cdots P(w_n|w_1^{n-1}) \quad \text{Chain}$$

$$\boxed{P(w_1)} \cdot P(w_2|w_1) \cdot P(w_3|w_2) \cdots P(w_n|w_{n-1}) \quad \text{Markov}$$

$$P(w_1 \boxed{w_0}) \cdot P(w_2|w_1) \cdot P(w_3|w_2) \cdots P(w_n|w_{n-1})$$

Likelihood of a **sequence of words**

$$P(w_1^n) = \prod_{k=1}^n P(w_k|w_{k-1})$$

Log Likelihood

Benefit?

$$\boxed{\log} P(w_1^n) = \boxed{\log} \prod_{k=1}^n P(w_k|w_{k-1}) = \sum_{k=1}^n \boxed{\log} P(w_k|w_{k-1})$$



Word Segmentation

- Word segmentation
 - Segment a chunk of string into a **sequence of words**.
 - Are there more than one possible sequence?
 - Choose the sequence that **most likely** appears in context.

youknow

$$P(you) \cdot P(know|you) > P(yo) \cdot P(uk|yo) \cdot P(now|uk)$$

$$\log(P(you) \cdot P(know|you)) > \log(P(yo) \cdot P(uk|yo) \cdot P(now|uk))$$

$$\log(P(you)) + \log(P(know|you)) > \log(P(yo)) + \log(P(uk|yo)) + \log(P(now|uk))$$

Perplexity

How to **evaluate** a **language model**?

$$\text{PL}(W) = P(w_1^n)^{-\frac{1}{n}} = \sqrt[n]{\frac{1}{P(w_1^n)}}$$

Inverse probability

Why?

$$\sqrt[n]{\prod_{k=1}^n \frac{1}{P(w_k | w_{k-1})}}$$

Branching factor?



Entropy

Entropy of word sequences whose length is n .

Measure of
information

$$W_n = \{w_1, w_2, \dots, w_n\}$$

$$H(W_n) = - \sum_{W'_n \in \mathcal{L}} p(W'_n) \log p(W'_n) \approx -\frac{1}{n} \log p(W'_n)$$

Entropy of language.

Shannon-
McMillan-
Breiman
Theorem

$$H(\mathcal{L}) = \lim_{n \rightarrow \infty} \frac{1}{n} H(W_n)$$

$$= \lim_{n \rightarrow \infty} -\frac{1}{n} \sum_{W'_n \in \mathcal{L}} p(W'_n) \log p(W'_n)$$

$$= \lim_{n \rightarrow \infty} -\frac{1}{n} \log p(W'_n)$$

Entropy vs. Perplexity

$$\boxed{H(W_n)} = -\frac{1}{n} \log P(w_1^n)$$

$$2^{H(W_n)} = 2^{-\frac{1}{n} \log P(w_1^n)}$$

$$= 2^{\log P(w_1^n) - \frac{1}{n}}$$

$$= P(w_1^n)^{-\frac{1}{n}} = \boxed{\text{PL}(W_n)}$$

Laplace Smoothing

$$P(x_1^n) \approx P(x_1) \cdot P(x_2|x_1) \cdot P(x_3|x_2) \cdots P(x_n|x_{n-1})$$

What if $P(x_1) = 0$? $P(x_1^n) \approx 0 \leftarrow \text{BAD!!}$

$$P(x_i) = \frac{C(x_i)}{\sum_k C(x_k)}$$

Laplace Smoothing

$$P_l(x_i) = \frac{C(x_i) + \alpha}{\sum_k (C(x_k) + \alpha)} = \frac{C(x_i) + \alpha}{\sum_k C(x_k) + \alpha|X|} = \frac{C(x_i) + \alpha}{N + \alpha|X|}$$

$$P_l(x_?) = \frac{C(x_?) + \alpha}{\sum_k C(x_k) + \alpha|X|} = \frac{\alpha}{N + \alpha|X|}$$



Laplace Smoothing

$$P(x_j|x_i) = \frac{C(x_i, x_j)}{\sum_k C(x_i, x_k)} = \frac{C(x_i, x_j)}{C(x_i)}$$

Laplace Smoothing

$$\begin{aligned} P_l(x_j|x_i) &= \frac{C(x_i, x_j) + \alpha}{\sum_k (C(x_i, x_k) + \alpha)} \\ &= \frac{C(x_i, x_j) + \alpha}{\sum_k C(x_i, x_k) + \alpha|X_{i,*}|} \\ &= \frac{C(x_i, x_j) + \alpha}{C(x_i) + \alpha|X_{i,*}|} \end{aligned}$$

$$P_l(x_?|x_i) = \frac{\alpha}{C(x_i) + \alpha|X_{i,*}|}$$



Discount Smoothing

- Issues with Laplace smoothing

- Unfair discounts.

$$\frac{1}{100} = 0.01 \rightarrow \frac{1 + 1}{100 + 10} = 0.018 \quad +0.008$$

$$\frac{10}{100} = 0.1 \rightarrow \frac{10 + 1}{100 + 10} = 0.1 \quad 0$$

$$\frac{50}{100} = 0.5 \rightarrow \frac{50 + 1}{100 + 10} = 0.46 \quad -0.04$$

- Unseen likelihood may get penalized too harshly when the minimum count is much greater than α .
 - How to reduce the gap between the **minimum** count and **unseen** count?

Discount Smoothing

Laplace

$$P_l(x?) = \frac{\alpha}{N + \alpha|X|}$$

$$P_l(x_i) = \frac{C(x_i) + \alpha}{N + \alpha|X|}$$

$$P_l(x?|x_i) = \frac{\alpha}{C(x_i) + \alpha|X_{i,*}|}$$

$$P_l(x_j|x_i) = \frac{C(x_i, x_j) + \alpha}{C(x_i) + \alpha|X_{i,*}|}$$

Discount

$$P_d(x?) = \alpha \cdot \min_k P(x_k)$$

$$P_d(x_i) = \frac{C(x_i) - P_d(x?)}{N}$$

$$P_d(x?|x_i) = \alpha \cdot \min_k P(x_k|x_i)$$

$$P_d(x_j|x_i) = \frac{C(x_i, x_j) - P_d(x?|x_i)}{C(x_i)}$$

Backoff

- Backoff
 - Bigrams are more **accurate** than unigrams.
 - Bigrams are **sparser** than unigrams.
 - Use bigrams in **general**, and use unigrams only bigrams **don't exist**.

$$P_b(x_j|x_i) = \begin{cases} P_{l|d}(x_j|x_i) & P(x_j|x_i) > 0 \\ \beta \cdot P_{l|d}(x_j) & \text{Otherwise} \end{cases}$$

How to measure?

$$\beta = \alpha \cdot \frac{\langle (P(x_j|x_i)) \rangle_{i,j}}{\langle (P(x_j)) \rangle_j}$$

Interpolation

- Interpolation
 - Unigrams and bigrams provide **different** but **useful** information .
 - Use them **both** with different **weights**.

$$\hat{P}(x_j|x_i) = \lambda_1 \cdot P_{l|d}(x_j) + \lambda_2 \cdot P_{l|d}(x_j|x_i)$$

$$\sum_k \lambda_k = 1$$